

```
In [1]: import pandas as pd
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix, recall_score, accuracy_score, precision_score

RANDOM_SEED = 2021
TEST_PCT = 0.3
LABELS = ["Normal", "Fraud"]
```

```
In [2]: dataset = pd.read_csv("C:/Users/Windows 10/Desktop/creditcard.csv")
```


```
In [3]: print("Any nulls in the dataset", dataset.isnull().values.any())
print('-----')
print("No. of unique labels", len(dataset['Class'].unique()))
print("Label values", dataset.Class.unique())

print('-----')
print("Break down of Normal and Fraud Transactions")
print(pd.value_counts(dataset['Class'], sort=True))
```

```
Any nulls in the dataset False
-----
No. of unique labels 2
Label values [0 1]
-----
Break down of Normal and Fraud Transactions
Class
0    284315
1      492
Name: count, dtype: int64
```


```
In [4]: count_classes = pd.value_counts(dataset['Class'], sort=True)
count_classes.plot(kind='bar', rot=0)
plt.xticks(range(len(dataset['Class'].unique())), dataset.Class.unique())
plt.title("Frequency by observation number")
plt.xlabel("Class")
plt.ylabel("Number of Observations")
```

```
Out[4]: Text(0, 0.5, 'Number of Observations')
```



```
In [5]: normal_dataset = dataset[dataset.Class == 0]
        fraud_dataset = dataset[dataset.Class == 1]

        bins = np.linspace(200,2500,100)
        plt.hist(normal_dataset.Amount,bins=bins,alpha=1,density=True,label='Normal')
        plt.hist(fraud_dataset.Amount,bins=bins,alpha=0.5,density=True,label='Fraud')
        plt.legend(loc='upper right')
        plt.title("Transcation Amount vs Percentage of Transcations")
        plt.xlabel("Transcation Amount (USD)")
        plt.ylabel("Percentage of Transcations")
        plt.show()
```



In [6]: dataset

Out[6]:

	Time	V1	V2	V3	V4	V5	V6	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.2395
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.0788
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.7914
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.2376
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.5925
...
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.9182
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.0243
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.2968
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.6867
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.5770

284807 rows × 31 columns

```
In [7]: sc = StandardScaler()
dataset['Time'] = sc.fit_transform(dataset['Time'].values.reshape(-1,1))
dataset['Amount'] = sc.fit_transform(dataset['Amount'].values.reshape(-1,1))
```

```
In [8]: raw_data = dataset.values
labels = raw_data[:, -1]
data = raw_data[:, 0:-1]
train_data, test_data, train_labels, test_labels = train_test_split(data, labels, test_s
```

```
In [9]: min_val = tf.reduce_min(train_data)
max_val = tf.reduce_max(train_data)

train_data = (train_data - min_val) / (max_val - min_val)
test_data = (test_data - min_val) / (max_val - min_val)

train_data = tf.cast(train_data, tf.float32)
test_data = tf.cast(test_data, tf.float32)
```

```
In [10]: train_labels = train_labels.astype(bool)
test_labels = test_labels.astype(bool)

normal_train_data = train_data[~train_labels]
normal_test_data = test_data[~test_labels]

fraud_train_data = train_data[train_labels]
fraud_test_data = test_data[test_labels]
print("No. of records in Fraud Train Data=", len(fraud_train_data))
print("No. of records in Normal Train Data=", len(normal_train_data))
print("No. of records in Fraud Test Data=", len(fraud_test_data))
print("No. of records in Normal Test Data=", len(normal_test_data))
```

```
No. of records in Fraud Train Data= 389
No. of records in Normal Train Data= 227456
No. of records in Fraud Test Data= 103
No. of records in Normal Test Data= 56859
```

```
In [11]: nb_epoch = 50
batch_size = 64
input_dim = normal_train_data.shape[1]
encoding_dim = 14
hidden_dim1 = int(encoding_dim / 2)
hidden_dim2 = 4
learning_rate = 1e-7
```

```
In [12]: input_layer = tf.keras.layers.Input(shape=(input_dim,))

encoder = tf.keras.layers.Dense(encoding_dim, activation="tanh", activity_regularizer
encoder = tf.keras.layers.Dropout(0.2)(encoder)
encoder = tf.keras.layers.Dense(hidden_dim1, activation='relu')(encoder)
encoder = tf.keras.layers.Dense(hidden_dim2, activation=tf.nn.leaky_relu)(encoder)

decoder = tf.keras.layers.Dense(hidden_dim1, activation='relu')(encoder)
decoder = tf.keras.layers.Dropout(0.2)(decoder)
decoder = tf.keras.layers.Dense(encoding_dim, activation='relu')(decoder)
decoder = tf.keras.layers.Dense(input_dim, activation='tanh')(decoder)
```

```
autoencoder = tf.keras.Model(inputs = input_layer, outputs = decoder)
autoencoder.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 30)]	0
dense (Dense)	(None, 14)	434
dropout (Dropout)	(None, 14)	0
dense_1 (Dense)	(None, 7)	105
dense_2 (Dense)	(None, 4)	32
dense_3 (Dense)	(None, 7)	35
dropout_1 (Dropout)	(None, 7)	0
dense_4 (Dense)	(None, 14)	112
dense_5 (Dense)	(None, 30)	450
=====		
Total params: 1168 (4.56 KB)		
Trainable params: 1168 (4.56 KB)		
Non-trainable params: 0 (0.00 Byte)		

```
In [13]: cp = tf.keras.callbacks.ModelCheckpoint(filepath="autoencoder_fraud.h5", mode='min',
early_stop = tf.keras.callbacks.EarlyStopping(
    monitor='val_loss',
    min_delta=0.0001,
    patience=10,
    verbose=11,
    mode='min',
    restore_best_weights=True
)
```

```
In [14]: autoencoder.compile(metrics=['accuracy'], loss='mean_squared_error', optimizer='adam'
```

```
In [15]: history = autoencoder.fit(normal_train_data, normal_train_data, epochs = nb_epoch,
    batch_size = batch_size, shuffle = True,
    validation_data = (test_data, test_data),
    verbose=1,
    callbacks = [cp, early_stop]).history
```

```
Epoch 1/50
3524/3554 [=====>.] - ETA: 0s - loss: 0.0050 - accuracy: 0.0525
Epoch 1: val_loss improved from inf to 0.00002, saving model to autoencoder_fraud.h5
3554/3554 [=====] - 7s 2ms/step - loss: 0.0050 - accuracy: 0.0524 - val_loss: 1.9844e-05 - val_accuracy: 0.0371
Epoch 2/50
 62/3554 [.....] - ETA: 5s - loss: 1.8694e-05 - accuracy: 0.0481


C:\Users\Windows 10\AppData\Roaming\Python\Python310\site-packages\keras\src\engine\training.py:3000: UserWarning: You are saving your model as an HDF5 file via `model.save()`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')`.
  saving_api.save_model(
```

```
3526/3554 [=====>.] - ETA: 0s - loss: 1.9268e-05 - accuracy:
0.0631
Epoch 2: val_loss did not improve from 0.00002
3554/3554 [=====] - 7s 2ms/step - loss: 1.9264e-05 - accuracy: 0.0633 - val_loss: 2.0042e-05 - val_accuracy: 0.1279
Epoch 3/50
3532/3554 [=====>.] - ETA: 0s - loss: 1.9425e-05 - accuracy:
0.0623
Epoch 3: val_loss did not improve from 0.00002
3554/3554 [=====] - 6s 2ms/step - loss: 1.9447e-05 - accuracy: 0.0624 - val_loss: 2.0213e-05 - val_accuracy: 0.0420
Epoch 4/50
3535/3554 [=====>.] - ETA: 0s - loss: 1.9589e-05 - accuracy:
0.0600
Epoch 4: val_loss did not improve from 0.00002
3554/3554 [=====] - 6s 2ms/step - loss: 1.9584e-05 - accuracy: 0.0601 - val_loss: 2.0390e-05 - val_accuracy: 0.0661
Epoch 5/50
3547/3554 [=====>.] - ETA: 0s - loss: 1.9570e-05 - accuracy:
0.0597
Epoch 5: val_loss did not improve from 0.00002
3554/3554 [=====] - 6s 2ms/step - loss: 1.9571e-05 - accuracy: 0.0598 - val_loss: 2.0407e-05 - val_accuracy: 0.1279
Epoch 6/50
3535/3554 [=====>.] - ETA: 0s - loss: 1.9551e-05 - accuracy:
0.0590
Epoch 6: val_loss did not improve from 0.00002
3554/3554 [=====] - 9s 2ms/step - loss: 1.9548e-05 - accuracy: 0.0591 - val_loss: 2.0087e-05 - val_accuracy: 0.0420
Epoch 7/50
3534/3554 [=====>.] - ETA: 0s - loss: 1.9536e-05 - accuracy:
0.0613
Epoch 7: val_loss did not improve from 0.00002
3554/3554 [=====] - 8s 2ms/step - loss: 1.9542e-05 - accuracy: 0.0612 - val_loss: 2.0120e-05 - val_accuracy: 0.0078
Epoch 8/50
3549/3554 [=====>.] - ETA: 0s - loss: 1.9550e-05 - accuracy:
0.0589
Epoch 8: val_loss did not improve from 0.00002
3554/3554 [=====] - 7s 2ms/step - loss: 1.9546e-05 - accuracy: 0.0590 - val_loss: 2.0156e-05 - val_accuracy: 0.0371
Epoch 9/50
3550/3554 [=====>.] - ETA: 0s - loss: 1.9536e-05 - accuracy:
0.0621
Epoch 9: val_loss did not improve from 0.00002
3554/3554 [=====] - 7s 2ms/step - loss: 1.9532e-05 - accuracy: 0.0623 - val_loss: 2.0003e-05 - val_accuracy: 0.2168
Epoch 10/50
3554/3554 [=====] - ETA: 0s - loss: 1.9525e-05 - accuracy:
0.0566
Epoch 10: val_loss did not improve from 0.00002
3554/3554 [=====] - 7s 2ms/step - loss: 1.9525e-05 - accuracy: 0.0566 - val_loss: 2.0068e-05 - val_accuracy: 0.2168
Epoch 11/50
3527/3554 [=====>.] - ETA: 0s - loss: 1.9535e-05 - accuracy:
0.0589
```

Epoch 11: val_loss did not improve from 0.00002
 Restoring model weights from the end of the best epoch: 1.
 3554/3554 [=====] - 7s 2ms/step - loss: 1.9524e-05 - accuracy: 0.0589 - val_loss: 2.0277e-05 - val_accuracy: 0.0814
 Epoch 11: early stopping

```
In [16]: plt.plot(history['loss'],linewidth = 2,label = 'Train')
plt.plot(history['val_loss'],linewidth = 2,label = 'Test')
plt.legend(loc='upper right')
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')

plt.show()
```



```
In [17]: test_x_predictions = autoencoder.predict(test_data)
mse = np.mean(np.power(test_data - test_x_predictions, 2),axis = 1)
error_df = pd.DataFrame({'Reconstruction_error':mse,
                        'True_class':test_labels})
```

1781/1781 [=====] - 2s 1ms/step

```
In [18]: threshold_fixed = 50
groups = error_df.groupby('True_class')
fig,ax = plt.subplots()


for name,group in groups:
    ax.plot(group.index,group.Reconstruction_error,marker='o',ms = 3.5,linestyle=
            label = "Fraud" if name==1 else "Normal")
```



```

ax.hlines(threshold_fixed,ax.get_xlim()[0],ax.get_xlim()[1],colors="r",zorder=100,1
ax.legend()
plt.title("Reconstructions error for normal and fraud data")
plt.ylabel("Reconstruction error")
plt.xlabel("Data point index")
plt.show()

```




```

In [19]: threshold_fixed = 52
pred_y = [1 if e > threshold_fixed else 0
          for e in
            error_df.Reconstruction_error.values]
error_df['pred'] = pred_y
conf_matrix = confusion_matrix(error_df.True_class,pred_y)

plt.figure(figsize = (4,4))
sns.heatmap(conf_matrix,xticklabels = LABELS,yticklabels = LABELS,annot = True,fmt=
plt.title("Confusion matrix")
plt.ylabel("True class")
plt.xlabel("Predicted class")
plt.show()

print("Accuracy :",accuracy_score(error_df['True_class'],error_df['pred']))
print("Recall :",recall_score(error_df['True_class'],error_df['pred']))
print("Precision :",precision_score(error_df['True_class'],error_df['pred']))

```



Accuracy : 0.9981917769741231

Recall : 0.0

Precision : 0.0

C:\Users\Windows 10\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\metrics_classification.py:1469: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))